

Skin Cancer Classification - An Educational Guide

In this tutorial we aim to provide a simple step-by-step guide to anyone who wants to work on the problem of skin lesion classification regardless of their level or expertise; from medical doctors, to master students and more experienced researchers.

Using this guide you will learn:

- How to load the data, visualise it and uncover more about the class distribution and meta-data.
- How to utilise architectures with varying complexity from a few convolutional layers to hundreds of them.
- How to train a model with appropriate optimisers and loss functions.
- How to rigorously test your trained model, providing not only metrics such as accuracy but also visualisations like confusion matrix and Grad — Cam.
- How to analyse and understand your results.

To conclude with, we will provide a few more tips that are usually utilised by the participants of the ISIC Challenges, that will help you increase your model's performance even more so that you can beat our performance and explore more advanced training schemes.

```
%pip install imageio
%pip install scikit-image
%pip install torch torchvision
%pip install numpy pandas scikit-learn scipy seaborn pillow matplotlib
```

```
Defaulting to user installation because normal site-packages is not
writeable
```

Collecting imageio

Downloading imageio-2.34.0-py3-none-any.whl (313 kB)

313.4/313.4 KB 963.7 kB/s eta

0:00:0000:0100:01

```
ent already satisfied: numpy in /usr/local/lib/python3.10/dist-  
packages (from imageio) (1.26.0)
```

```
Requirement already satisfied: pillow>=8.3.2 in /usr/lib/python3/dist-packages (from imageio) (9.0.1)
```

```
Installing collected packages: imageio
```

Successfully installed imageio-2.34.0

```
Defaulting to user installation because normal site-packages is not
writable
```

Collecting scikit-image

```
Downloading scikit image-0.22.0-cp310-cp310-
```

manylinux 2 17 x86_64.manylinux2014 x86_64.whl (14.7 MB)

14.7/14.7 MB 13.9 MB/s eta

```

0:00:0000:0100:01
ent already satisfied: pillow>=9.0.1 in /usr/lib/python3/dist-packages
(from scikit-image) (9.0.1)
Requirement already satisfied: networkx>=2.8 in
/usr/local/lib/python3.10/dist-packages (from scikit-image) (3.1)
Collecting lazy_loader>=0.3
  Downloading lazy_loader-0.3-py3-none-any.whl (9.1 kB)
Requirement already satisfied: numpy>=1.22 in
/usr/local/lib/python3.10/dist-packages (from scikit-image) (1.26.0)
Requirement already satisfied: packaging>=21 in
/usr/local/lib/python3.10/dist-packages (from scikit-image) (23.2)
Requirement already satisfied: imageio>=2.27 in
/usr/cs/grad/masters/2024/akalapal/.local/lib/python3.10/site-packages
(from scikit-image) (2.34.0)
Requirement already satisfied: scipy>=1.8 in
/usr/local/lib/python3.10/dist-packages (from scikit-image) (1.11.3)
Collecting tifffile>=2022.8.12
  Downloading tifffile-2024.2.12-py3-none-any.whl (224 kB)
224.5/224.5 KB 1.1 MB/s eta
0:00:00a 0:00:01
age
Successfully installed lazy_loader-0.3 scikit-image-0.22.0 tifffile-
2024.2.12
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: torch in
/usr/local/lib/python3.10/dist-packages (2.1.0)
Requirement already satisfied: torchvision in
/usr/local/lib/python3.10/dist-packages (0.16.0)
Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26 in
/usr/local/lib/python3.10/dist-packages (from torch) (8.9.2.26)
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in
/usr/local/lib/python3.10/dist-packages (from torch) (10.3.2.106)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch) (12.1.105)
Requirement already satisfied: sympy in
/usr/local/lib/python3.10/dist-packages (from torch) (1.12)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105
in /usr/local/lib/python3.10/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in
/usr/local/lib/python3.10/dist-packages (from torch) (11.0.2.54)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch) (12.1.105)
Requirement already satisfied: jinja2 in /usr/lib/python3/dist-
packages (from torch) (3.0.3)
Requirement already satisfied: networkx in
/usr/local/lib/python3.10/dist-packages (from torch) (3.1)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in
/usr/local/lib/python3.10/dist-packages (from torch) (11.4.5.107)

```

Requirement already satisfied: nvidia-nccl-cu12==2.18.1 in
/usr/local/lib/python3.10/dist-packages (from torch) (2.18.1)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch) (12.1.105)
Requirement already satisfied: triton==2.1.0 in
/usr/local/lib/python3.10/dist-packages (from torch) (2.1.0)
Requirement already satisfied: filelock in
/usr/local/lib/python3.10/dist-packages (from torch) (3.12.4)
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in
/usr/local/lib/python3.10/dist-packages (from torch) (12.1.3.1)
Requirement already satisfied: fsspec in
/usr/local/lib/python3.10/dist-packages (from torch) (2023.6.0)
Requirement already satisfied: nvidia-cusparse-cu12==12.1.0.106 in
/usr/local/lib/python3.10/dist-packages (from torch) (12.1.0.106)
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.10/dist-packages (from torch) (4.8.0)
Requirement already satisfied: nvidia-nvjitlink-cu12 in
/usr/local/lib/python3.10/dist-packages (from nvidia-cusolver-
cu12==11.4.5.107->torch) (12.2.140)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/lib/python3/dist-packages (from torchvision) (9.0.1)
Requirement already satisfied: requests in
/usr/local/lib/python3.10/dist-packages (from torchvision) (2.31.0)
Requirement already satisfied: numpy in
/usr/local/lib/python3.10/dist-packages (from torchvision) (1.26.0)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests->torchvision)
(2.2.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/lib/python3/dist-
packages (from requests->torchvision) (3.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->torchvision)
(3.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/lib/python3/dist-packages (from requests->torchvision)
(2020.6.20)
Requirement already satisfied: mpmath>=0.19 in
/usr/local/lib/python3.10/dist-packages (from sympy->torch) (1.3.0)
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: numpy in
/usr/local/lib/python3.10/dist-packages (1.26.0)
Requirement already satisfied: pandas in
/usr/local/lib/python3.10/dist-packages (2.1.1)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.10/dist-packages (1.3.1)
Requirement already satisfied: scipy in
/usr/local/lib/python3.10/dist-packages (1.11.3)
Collecting seaborn

Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)

294.9/294.9 KB 1.2 MB/s eta

0:00:00a 0:00:01

Requirement already satisfied: pillow in /usr/lib/python3/dist-packages (9.0.1)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.8.0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)

Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.3)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.2.0)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.43.1)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.1.1)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/lib/python3/dist-packages (from matplotlib) (2.4.7)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (23.2)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)

Requirement already satisfied: six>=1.5 in /usr/lib/python3/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)

Installing collected packages: seaborn

Successfully installed seaborn-0.13.2

```
import torch
from torch import nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
import numpy as np
import os
import shutil
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix,
precision_recall_fscore_support
import scipy.ndimage
from scipy import misc
```

```

from glob import glob
from scipy import stats
from sklearn.preprocessing import LabelEncoder, StandardScaler
import skimage
import imageio
import seaborn as sns
from PIL import Image
import glob
import matplotlib.pyplot as plt
import matplotlib
%matplotlib inline

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device

device(type='cuda')

```

What about data?

The HAM10000 ("Human Against Machine with 10000 training images") dataset which contains 10,015 dermatoscopic images was made publically available by the Harvard database on June 2018 in the hopes to provide training data for automating the process of skin cancer lesion classifications. The motivation behind this act was to provide the public with an abundance and variability of data source for machine learning training purposes such that the results may be compared with that of human experts. If successful, the applications would bring cost and time saving regimes to hospitals and medical professions alike.

Apart from the 10,015 images, a metadata file with demographic information of each lesion is provided as well. More than 50% of lesions are confirmed through histopathology (histo), the ground truth for the rest of the cases is either follow-up examination (follow_up), expert consensus (consensus), or confirmation by in-vivo confocal microscopy (confocal)

You can download the dataset here: <https://dataverse.harvard.edu/dataset.xhtml?persistentId=doi:10.7910/DVN/DBW86T>

The 7 classes of skin cancer lesions included in this dataset are:

1. Melanocytic nevi
2. Melanoma
3. Benign keratosis-like lesions
4. Basal cell carcinoma
5. Actinic keratoses
6. Vascular lesions
7. Dermatofibroma

Let's analyze the metadata of the dataset

I have updated my path of data_dir.

```

# importing metadata and checking for its shape
data_dir = "/usr/cs/grad/masters/2024/akalapa/Desktop/hw4/HAM10000"

metadata = pd.read_csv(data_dir + '/HAM10000_metadata')
print(metadata.shape)

# label encoding the seven classes for skin cancers

le = LabelEncoder()
le.fit(metadata['dx'])
LabelEncoder()
print("Classes:", list(le.classes_))

metadata['label'] = le.transform(metadata["dx"])
metadata.sample(10)

(10015, 8)
Classes: ['akiec', 'bcc', 'bkl', 'df', 'mel', 'nv', 'vasc']

```

	lesion_id	image_id	dx	dx_type	age	sex	\
4242	HAM_0001256	ISIC_0025997	nv	follow_up	50.0	female	
2505	HAM_0006800	ISIC_0030778	bcc	histo	85.0	female	
5822	HAM_0001925	ISIC_0027583	nv	follow_up	45.0	male	
5655	HAM_0000438	ISIC_0025532	nv	follow_up	45.0	female	
3661	HAM_0006375	ISIC_0026114	nv	follow_up	50.0	female	
5419	HAM_0004493	ISIC_0030325	nv	follow_up	60.0	male	
7675	HAM_0007265	ISIC_0033639	nv	histo	20.0	male	
8749	HAM_0000896	ISIC_0025051	nv	histo	45.0	female	
2659	HAM_0002397	ISIC_0029192	bcc	histo	85.0	female	
3418	HAM_0004328	ISIC_0027164	nv	follow_up	45.0	female	

	localization	dataset	label
4242	trunk	vidir_molemax	5
2505	face	vidir_modern	1
5822	trunk	vidir_molemax	5
5655	lower extremity	vidir_molemax	5
3661	back	vidir_molemax	5
5419	upper extremity	vidir_molemax	5
7675	chest	vidir_modern	5
8749	lower extremity	rosendahl	5
2659	neck	vidir_modern	1
3418	abdomen	vidir_molemax	5

```

# Getting a sense of what the distribution of each column looks like

fig = plt.figure(figsize=(40,25))

ax1 = fig.add_subplot(221)
metadata['dx'].value_counts().plot(kind='bar', ax=ax1)
ax1.set_ylabel('Count', size=50)
ax1.set_title('Cell Type', size = 50)

```

```
ax2 = fig.add_subplot(222)
metadata['sex'].value_counts().plot(kind='bar', ax=ax2)
ax2.set_ylabel('Count', size=50)
ax2.set_title('Sex', size=50);

ax3 = fig.add_subplot(223)
metadata['localization'].value_counts().plot(kind='bar')
ax3.set_ylabel('Count', size=50)
ax3.set_title('Localization', size=50)

ax4 = fig.add_subplot(224)
sample_age = metadata[pd.notnull(metadata['age'])]
sns.distplot(sample_age['age'], fit=stats.norm, color='red');
ax4.set_title('Age', size = 50)
ax4.set_xlabel('Year', size=50)

plt.tight_layout()
plt.show()
```

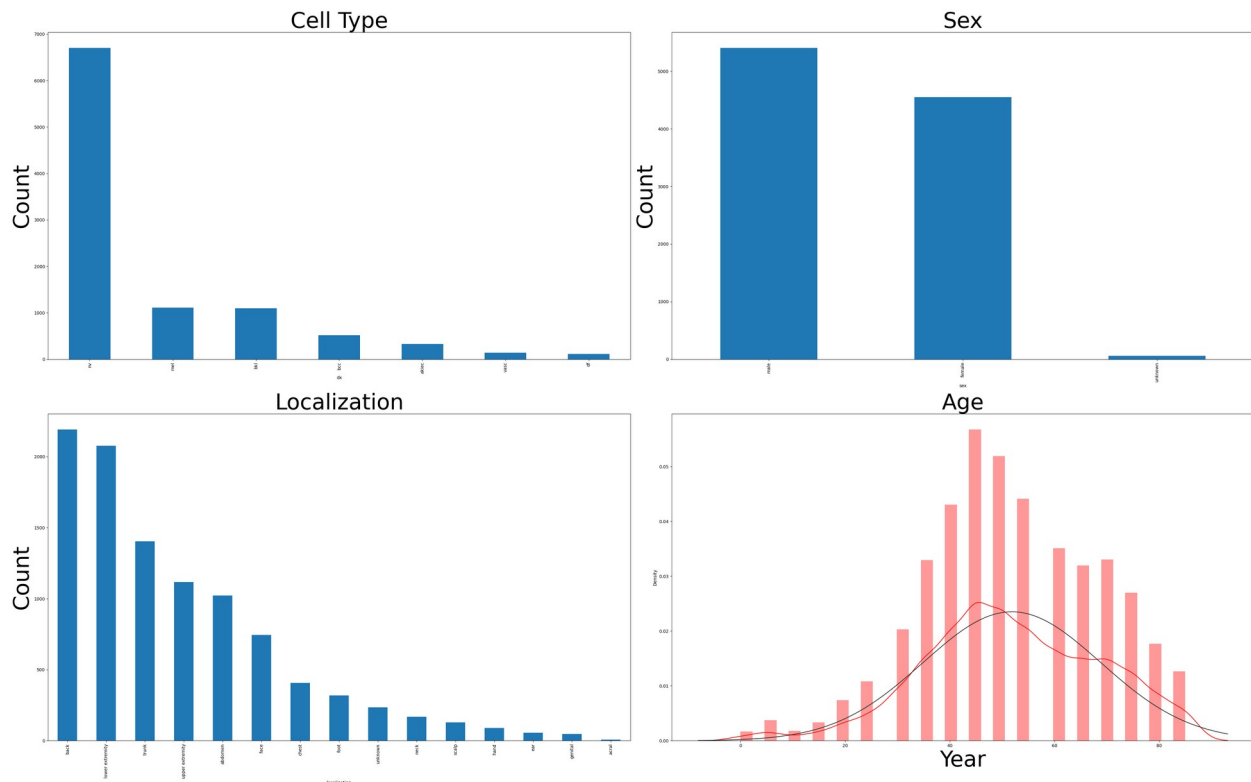
/tmp/ipykernel_1929082/965783653.py:23: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(sample_age['age'], fit=stats.norm, color='red');
```



As you can see there is imbalance in the number of images per class. There are much more images for the lesion type "Melanocytic Nevi" compared to other types. This is an usual occurrence for medical datasets and so it is very important to analyze the data from beforehand.

Let's visualize some examples

I removed the part ("`/`" + `label.values[0]` +) from `im_sample` to read the images according to their labels and classes.

```
#Visualizing the images

label = [ 'akiec', 'bcc', 'bkl', 'df', 'mel', 'nv', 'vasc' ]
label_images = []
classes = [ 'actinic keratoses', 'basal cell carcinoma', 'benign
keratosis-like lesions',
            'dermatofibroma', 'melanoma', 'melanocytic nevi', 'vascular
lesions' ]

fig = plt.figure(figsize=(55, 55))
k = range(7)

for i in label:
    sample = metadata[metadata['dx'] == i]['image_id'][:5]
    label_images.extend(sample)

for position, ID in enumerate(label_images):
    lbl = metadata[metadata['image_id'] == ID]['dx']
```



```
im_sample = data_dir + f'/{ID}.jpg'
im_sample = imageio.imread(im_sample)

plt.subplot(7,5,position+1)
plt.imshow(im_sample)
plt.axis('off')





if position%5 == 0:
    title = int(position/5)
    plt.title(classes[title], loc='left', size=50, weight="bold")
```

```
plt.tight_layout()
plt.show()
```

/tmp/ipykernel_1929082/3944126756.py:18: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `io.v3.imread`. To keep the current behavior (and make this warning disappear) use `import imageio.v2 as imageio` or call `imageio.v2.imread` directly.

```
im_sample = imageio.imread(im_sample)
```






actinic keratoses




basal cell carcinoma






benign keratosis-like lesions




dermatofibroma






melanoma



melanocytic nevi



vascular lesions



Median Frequency Balancing

As we saw above that there is class imbalance in our dataset. To solve that we use this method.

module 'numpy' has no attribute 'float'. `np.float` was a deprecated alias for the builtin `float`. To avoid this error in existing code, i removed the part `np.` for `class_weights` and `counts`.

```
import numpy as np
import pandas as pd

print(metadata['dx'].value_counts())
```

```

print(metadata[metadata['dx']=='nv']['dx'].value_counts())
label = ['akiec', 'bcc', 'bkl', 'df', 'mel', 'nv', 'vasc']

def estimate_weights_mfb(label):
    class_weights = np.zeros_like(label, dtype=float)
    counts = np.zeros_like(label)
    for i, l in enumerate(label):
        counts[i] = metadata[metadata['dx'] == str(l)]
        ['dx'].value_counts().iloc[0]
    counts = counts.astype(float)
    median_freq = np.median(counts)
    for i, label in enumerate(label):
        class_weights[i] = median_freq / counts[i]
    return class_weights

classweight = estimate_weights_mfb(label)

for i in range(len(label)):
    print(label[i], ":", classweight[i])

```

dx	
nv	6705
mel	1113
bkl	1099
bcc	514
akiec	327
vasc	142
df	115

Name: count, dtype: int64

dx	
nv	6705

Name: count, dtype: int64

akiec : 1.5718654434250765

bcc : 1.0

bkl : 0.467697907188353

df : 4.469565217391304

mel : 0.4618149146451033

nv : 0.07665920954511558

vasc : 3.619718309859155

Pre-processing the dataset

Before we load the data we need to alter the dataset structure. When you download the dataset, all the images are together in a folder. To use Pytorch dataloader we need to segregate the images into folders of their respective labels. You can use the following script to automate the process.

```

import os
import shutil

```

```

data_dir = "/usr/cs/grad/masters/2024/akalapal/Desktop/hw4/HAM10000/"
dest_dir = data_dir + "test/"
metadata = pd.read_csv(data_dir + '/HAM10000_metadata')

label = ['bkl', 'nv', 'df', 'mel', 'vasc', 'bcc', 'akiec']
label_images = []

for i in label:
    os.mkdir(dest_dir + str(i) + "/")
    sample = metadata[metadata['dx'] == i]['image_id']
    label_images.extend(sample)
    for id in label_images:
        shutil.copyfile((data_dir + id + ".jpg"), (dest_dir + i +
"/"+id+".jpg"))
    label_images=[]

```

Data Augmentation

It is a common fact that medical data is scarce. But to learn a very good model, the network needs a lot of data. So to tackle the problem we perform data augmentation.

First we normalize the images. Data normalization is an important step which ensures that each input parameter (pixel, in this case) has a similar data distribution. This makes convergence faster while training the network. Data normalization is done by subtracting the mean from each pixel and then dividing the result by the standard deviation. The distribution of such data would resemble a Gaussian curve centered at zero. Since, skin lesion images are natural images, we use the normalization values (mean and standard deviation) of Imagenet dataset.

We also perform data augmentation:

- Flipping the image horizontally: *RandomHorizontalFlip()*
- Rotating image 60 degrees: *RandomRotation()*. 60 degrees is chosen as best practice. You can experiment with other angles.

The augmentation is applied using the *transform.Compose()* function of Pytorch. Take note, we only augment the training set. This is because, augmentation is done to aid the training process. So there is no point in augmenting the test set.

I changed my path to my original dataset path.

```

data_dir =
"/usr/cs/grad/masters/2024/akalapal/Desktop/hw4/HAM10000/test"

# normalization values for pretrained resnet on Imagenet
norm_mean = (0.4914, 0.4822, 0.4465)
norm_std = (0.2023, 0.1994, 0.2010)

batch_size = 50
validation_batch_size = 10
test_batch_size = 10

```

```

# We compute the weights of individual classes and convert them to
tensors
class_weights = estimate_weights_mfb(label)
class_weights = torch.FloatTensor(class_weights)

transform_train = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(degrees=60),
    transforms.ToTensor(),
    transforms.Normalize(norm_mean, norm_std),
])

transform_test = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465),
(0.2023, 0.1994, 0.2010)),
])

```

Train, Test and Validation Split

We split the entire dataset into 3 parts:

- Train: 80%
- Test: 20%
- Validation: 16%

The splitting is done class wise so that we have equal representation of all classes in each subset of the data.

I have updated the line `def init(self, class_vector, test_size)` by adding `test_size` as an argument for the function.

```

import torch as th
import math

test_size = 0.2
val_size = 0.2
class Sampler(object):
    """Base class for all Samplers.
    """

    def __init__(self, data_source):
        pass

    def __iter__(self):
        raise NotImplementedError

```

```

def __len__(self):
    raise NotImplementedError

class StratifiedSampler(Sampler):
    """Stratified Sampling
    Provides equal representation of target classes
    """
    def __init__(self, class_vector, test_size):
        """
        Arguments
        -----
        class_vector : torch tensor
            a vector of class labels
        batch_size : integer
            batch_size
        """
        self.n_splits = 1
        self.class_vector = class_vector
        self.test_size = test_size

    def gen_sample_array(self):
        try:
            from sklearn.model_selection import StratifiedShuffleSplit
        except:
            print('Need scikit-learn for this functionality')
        import numpy as np

        s = StratifiedShuffleSplit(n_splits=self.n_splits,
test_size=self.test_size)
        X = th.randn(self.class_vector.size(0),2).numpy()
        y = self.class_vector.numpy()
        s.get_n_splits(X, y)

        train_index, test_index= next(s.split(X, y))
        return train_index, test_index

    def __iter__(self):
        return iter(self.gen_sample_array())

    def __len__(self):
        return len(self.class_vector)

dataset = torchvision.datasets.ImageFolder(root= data_dir)
data_label = [s[1] for s in dataset.samples]

ss = StratifiedSampler(torch.FloatTensor(data_label), test_size)
pre_train_indices, test_indices = ss.gen_sample_array()
# The "pre" is necessary to use array to identify train/ val indices
with indices generated by second sampler

```

```

train_label = np.delete(data_label, test_indices, None)
ss = StratifiedSampler(torch.FloatTensor(train_label), test_size)
train_indices, val_indices = ss.gen_sample_array()
indices = {'train': pre_train_indices[train_indices], # Indices of
second sampler are used on pre_train_indices
           'val': pre_train_indices[val_indices], # Indices of second
sampler are used on pre_train_indices
           'test': test_indices
          }

train_indices = indices['train']
val_indices = indices['val']
test_indices = indices['test']
print("Train Data Size:", len(train_indices))
print("Test Data Size:", len(test_indices))
print("Validation Data Size:", len(val_indices))

Train Data Size: 6409
Test Data Size: 2003
Validation Data Size: 1603

```

Now we use Pytorch data loader to load the dataset into the memory.

```

SubsetRandomSampler = torch.utils.data.sampler.SubsetRandomSampler

dataset = torchvision.datasets.ImageFolder(root= data_dir,
transform=transform_train)

train_samples = SubsetRandomSampler(train_indices)
val_samples = SubsetRandomSampler(val_indices)
test_samples = SubsetRandomSampler(test_indices)

train_data_loader = torch.utils.data.DataLoader(dataset,
batch_size=batch_size, shuffle=False, num_workers=1, sampler=
train_samples)
validation_data_loader = torch.utils.data.DataLoader(dataset,
batch_size=validation_batch_size, shuffle=False, sampler=val_samples)

dataset = torchvision.datasets.ImageFolder(root= data_dir,
transform=transform_test)
test_data_loader = torch.utils.data.DataLoader(dataset,
batch_size=test_batch_size, shuffle=False, sampler=test_samples)

```

Let us see some of the training images.

```

# functions to show an image
fig = plt.figure(figsize=(10, 15))
def imshow(img):

```




```
img = img / 2 + 0.5      # denormalize change this
npimg = img.numpy()
plt.imshow(np.transpose(npimg, (1, 2, 0)))
```

```
# get some random training images
dataiter = iter(train_data_loader)
images, labels = next(dataiter)
```

```
# show images
imshow(torchvision.utils.make_grid(images))
#classes = ['bkl', 'nv', 'df', 'mel', 'vasc', 'bcc', 'akiec']
# print labels
print(' '.join('%5s, ' % classes[labels[j]] for j in
range(len(labels))))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

benign keratosis-like lesions, melanoma, melanocytic nevi,
melanocytic nevi, melanocytic nevi, melanocytic nevi, benign
keratosis-like lesions, benign keratosis-like lesions, basal cell
carcinoma, actinic keratoses, melanocytic nevi, actinic keratoses,
melanocytic nevi, melanocytic nevi, melanocytic nevi, melanocytic
nevi, actinic keratoses, melanocytic nevi, melanoma, melanocytic
nevi, melanocytic nevi, basal cell carcinoma, melanocytic nevi,
basal cell carcinoma, melanocytic nevi, melanocytic nevi,
melanocytic nevi, benign keratosis-like lesions, basal cell
carcinoma, melanocytic nevi, melanocytic nevi, melanocytic nevi,
melanocytic nevi, melanocytic nevi, melanocytic nevi, melanocytic
nevi, melanocytic nevi, melanocytic nevi, melanocytic nevi,
melanocytic nevi, dermatofibroma, melanocytic nevi, melanocytic
nevi, melanocytic nevi, melanoma, melanocytic nevi, melanocytic
nevi, melanocytic nevi, melanocytic nevi, benign keratosis-like
lesions,



Define a Convolutional Neural Network

Pytorch makes it very easy to define a neural network. We have layers like Convolutions, ReLU non-linearity, Maxpooling etc. directly from torch library.

In this tutorial, we use The LeNet architecture introduced by LeCun et al. in their 1998 paper, [Gradient-Based Learning Applied to Document Recognition](#). As the name of the paper suggests, the authors' implementation of LeNet was used primarily for OCR and character recognition in documents.

The LeNet architecture is straightforward and small, (in terms of memory footprint), making it perfect for teaching the basics of CNNs.

```
num_classes = len(classes)
class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()
```

```

        self.conv1 = nn.Conv2d(3, 6, (5,5), padding=2)
        self.conv2 = nn.Conv2d(6, 16, (5,5))
        self.fc1    = nn.Linear(16*54*54, 120)
        self.fc2    = nn.Linear(120, 84)
        self.fc3    = nn.Linear(84, num_classes)
    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), (2,2))
        x = F.max_pool2d(F.relu(self.conv2(x)), (2,2))
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
    def num_flat_features(self, x):
        size = x.size()[1:]
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = LeNet()
net = net.to(device)

```

Define a Loss function and Optimizer

Let's use a Classification Cross-Entropy loss.

$$H_{y'}(y) := - \sum_i y_i' \log(y_i)$$

The most common and effective Optimizer currently used is **Adam: Adaptive Moments**. You can look [here](#) for more information.

```

import torch.optim as optim

class_weights = class_weights.to(device)
criterion = nn.CrossEntropyLoss(weight = class_weights)
optimizer = optim.Adam(net.parameters(), lr=1e-5)
print(net)

LeNet(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2))
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=46656, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=7, bias=True)
)

```

These are some helper functions to evaluate the training process.

```

from sklearn.metrics import accuracy_score

def get_accuracy(predicted, labels):
    batch_len, correct= 0, 0
    batch_len = labels.size(0)
    correct = (predicted == labels).sum().item()
    return batch_len, correct

def evaluate(model, val_loader):
    losses= 0
    num_samples_total=0
    correct_total=0
    model.eval()
    for inputs, labels in val_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        out = model(inputs)
        _, predicted = torch.max(out, 1)
        loss = criterion(out, labels)
        losses += loss.item()
        b_len, corr = get_accuracy(predicted, labels)
        num_samples_total +=b_len
        correct_total +=corr
    accuracy = correct_total/num_samples_total
    losses = losses/len(val_loader)
    return losses, accuracy

```

Train the network

This is when things start to get interesting. We simply loop over the training data iterator, and feed the inputs to the network and optimize.

```

# number of loops over the dataset
num_epochs = 50
accuracy = []
val_accuracy = []
losses = []
val_losses = []

for epoch in range(num_epochs):
    running_loss = 0.0
    correct_total= 0.0
    num_samples_total=0.0
    for i, data in enumerate(train_data_loader):
        # get the inputs
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        # set the parameter gradients to zero
        optimizer.zero_grad()

```

```

    # forward + backward + optimize
    outputs = net(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    #compute accuracy
    _, predicted = torch.max(outputs, 1)
    b_len, corr = get_accuracy(predicted, labels)
    num_samples_total += b_len
    correct_total += corr
    running_loss += loss.item()

running_loss /= len(train_data_loader)
train_accuracy = correct_total/num_samples_total
val_loss, val_acc = evaluate(net, validation_data_loader)

print('Epoch: %d' %(epoch+1))
print('Loss: %.3f Accuracy: %.3f' %(running_loss, train_accuracy))
print('Validation Loss: %.3f Val Accuracy: %.3f' %(val_loss,
val_acc))

losses.append(running_loss)
val_losses.append(val_loss)
accuracy.append(train_accuracy)
val_accuracy.append(val_acc)
print('Finished Training')

Epoch: 1
Loss: 1.340 Accuracy: 0.651
Validation Loss: 1.135 Val Accuracy: 0.659
Epoch: 2
Loss: 1.129 Accuracy: 0.639
Validation Loss: 1.055 Val Accuracy: 0.640
Epoch: 3
Loss: 1.077 Accuracy: 0.617
Validation Loss: 1.023 Val Accuracy: 0.651
Epoch: 4
Loss: 1.037 Accuracy: 0.604
Validation Loss: 0.997 Val Accuracy: 0.603
Epoch: 5
Loss: 1.008 Accuracy: 0.608
Validation Loss: 0.963 Val Accuracy: 0.593
Epoch: 6
Loss: 0.990 Accuracy: 0.610
Validation Loss: 0.945 Val Accuracy: 0.620
Epoch: 7
Loss: 0.976 Accuracy: 0.612

```

Validation Loss: 0.940 Val Accuracy: 0.614
Epoch: 8
Loss: 0.972 Accuracy:0.606
Validation Loss: 0.943 Val Accuracy: 0.619
Epoch: 9
Loss: 0.962 Accuracy:0.605
Validation Loss: 0.930 Val Accuracy: 0.600
Epoch: 10
Loss: 0.956 Accuracy:0.605
Validation Loss: 0.922 Val Accuracy: 0.623
Epoch: 11
Loss: 0.949 Accuracy:0.609
Validation Loss: 0.912 Val Accuracy: 0.619
Epoch: 12
Loss: 0.940 Accuracy:0.610
Validation Loss: 0.919 Val Accuracy: 0.617
Epoch: 13
Loss: 0.933 Accuracy:0.601
Validation Loss: 0.906 Val Accuracy: 0.623
Epoch: 14
Loss: 0.930 Accuracy:0.617
Validation Loss: 0.897 Val Accuracy: 0.639
Epoch: 15
Loss: 0.919 Accuracy:0.616
Validation Loss: 0.895 Val Accuracy: 0.627
Epoch: 16
Loss: 0.915 Accuracy:0.612
Validation Loss: 0.884 Val Accuracy: 0.628
Epoch: 17
Loss: 0.916 Accuracy:0.613
Validation Loss: 0.875 Val Accuracy: 0.615
Epoch: 18
Loss: 0.902 Accuracy:0.615
Validation Loss: 0.869 Val Accuracy: 0.621
Epoch: 19
Loss: 0.903 Accuracy:0.617
Validation Loss: 0.886 Val Accuracy: 0.618
Epoch: 20
Loss: 0.901 Accuracy:0.618
Validation Loss: 0.883 Val Accuracy: 0.621
Epoch: 21
Loss: 0.898 Accuracy:0.616
Validation Loss: 0.869 Val Accuracy: 0.623
Epoch: 22
Loss: 0.893 Accuracy:0.618
Validation Loss: 0.870 Val Accuracy: 0.636
Epoch: 23
Loss: 0.888 Accuracy:0.620
Validation Loss: 0.871 Val Accuracy: 0.644

Epoch: 24
Loss: 0.884 Accuracy:0.620
Validation Loss: 0.880 Val Accuracy: 0.609
Epoch: 25
Loss: 0.875 Accuracy:0.624
Validation Loss: 0.867 Val Accuracy: 0.627
Epoch: 26
Loss: 0.882 Accuracy:0.624
Validation Loss: 0.870 Val Accuracy: 0.638
Epoch: 27
Loss: 0.870 Accuracy:0.623
Validation Loss: 0.868 Val Accuracy: 0.616
Epoch: 28
Loss: 0.865 Accuracy:0.625
Validation Loss: 0.853 Val Accuracy: 0.629
Epoch: 29
Loss: 0.864 Accuracy:0.621
Validation Loss: 0.862 Val Accuracy: 0.654
Epoch: 30
Loss: 0.861 Accuracy:0.628
Validation Loss: 0.854 Val Accuracy: 0.652
Epoch: 31
Loss: 0.861 Accuracy:0.624
Validation Loss: 0.840 Val Accuracy: 0.634
Epoch: 32
Loss: 0.861 Accuracy:0.621
Validation Loss: 0.862 Val Accuracy: 0.596
Epoch: 33
Loss: 0.855 Accuracy:0.623
Validation Loss: 0.843 Val Accuracy: 0.641
Epoch: 34
Loss: 0.847 Accuracy:0.637
Validation Loss: 0.866 Val Accuracy: 0.589
Epoch: 35
Loss: 0.847 Accuracy:0.627
Validation Loss: 0.844 Val Accuracy: 0.634
Epoch: 36
Loss: 0.850 Accuracy:0.633
Validation Loss: 0.853 Val Accuracy: 0.623
Epoch: 37
Loss: 0.843 Accuracy:0.632
Validation Loss: 0.848 Val Accuracy: 0.622
Epoch: 38
Loss: 0.842 Accuracy:0.629
Validation Loss: 0.833 Val Accuracy: 0.651
Epoch: 39
Loss: 0.842 Accuracy:0.630
Validation Loss: 0.835 Val Accuracy: 0.634
Epoch: 40

```
Loss: 0.829 Accuracy:0.638
Validation Loss: 0.841 Val Accuracy: 0.623
Epoch: 41
Loss: 0.835 Accuracy:0.635
Validation Loss: 0.829 Val Accuracy: 0.641
Epoch: 42
Loss: 0.831 Accuracy:0.636
Validation Loss: 0.841 Val Accuracy: 0.645
Epoch: 43
Loss: 0.837 Accuracy:0.638
Validation Loss: 0.853 Val Accuracy: 0.605
Epoch: 44
Loss: 0.833 Accuracy:0.639
Validation Loss: 0.838 Val Accuracy: 0.633
Epoch: 45
Loss: 0.822 Accuracy:0.631
Validation Loss: 0.829 Val Accuracy: 0.635
Epoch: 46
Loss: 0.829 Accuracy:0.640
Validation Loss: 0.839 Val Accuracy: 0.651
Epoch: 47
Loss: 0.825 Accuracy:0.636
Validation Loss: 0.821 Val Accuracy: 0.635
Epoch: 48
Loss: 0.823 Accuracy:0.636
Validation Loss: 0.826 Val Accuracy: 0.666
Epoch: 49
Loss: 0.823 Accuracy:0.637
Validation Loss: 0.819 Val Accuracy: 0.651
Epoch: 50
Loss: 0.821 Accuracy:0.642
Validation Loss: 0.819 Val Accuracy: 0.653
Finished Training
```

Plot the training and validation loss curves.


```
#plt.plot(losses)
#plt.show()

epoch = range(1, num_epochs+1)
fig = plt.figure(figsize=(10, 15))
plt.subplot(2,1,2)
plt.plot(epoch, losses, label='Training loss')
plt.plot(epoch, val_losses, label='Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.legend()
plt.figure()
plt.show()
```

```

fig = plt.figure(figsize=(10, 15))
plt.subplot(2,1,2)
plt.plot(epoch, accuracy, label='Training accuracy')
plt.plot(epoch, val_accuracy, label='Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.legend()
plt.figure()
plt.show()

```



<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>

Test the network on the test data

We have trained the network over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs, and checking it against the ground-truth. If the prediction is correct, we add the sample to the list of correct predictions.


Okay, first step. Let us display an image from the test set to get familiar.

```
fig = plt.figure(figsize=(10, 15))
dataiter = iter(test_data_loader)
images, labels = next(dataiter)

# print images
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s, ' % classes[labels[j]] for j in
range(len(labels))))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

GroundTruth: benign keratosis-like lesions, melanocytic nevi, melanocytic nevi, melanocytic nevi, melanocytic nevi, benign keratosis-like lesions, actinic keratoses, benign keratosis-like lesions, melanocytic nevi, melanocytic nevi,



Okay, now let us check the performance on the test network:

```
correct = 0
total = 0
net.eval()
with torch.no_grad():
    for data in test_data_loader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the test images: 63 %

That looks better than chance, which is about 14% accuracy (randomly picking a class out of 7 classes). Seems like the network learnt something. But maybe it doesn't learn all the classes equally.

Let's check which classes that performed well, and which did not.

```
class_correct = list(0. for i in range(len(classes)))
class_total = list(1e-7 for i in range(len(classes)))
with torch.no_grad():
    for data in test_data_loader:
        images, labels = data
```

```

images, labels = images.to(device), labels.to(device)
outputs = net(images)
_, predicted = torch.max(outputs, 1)
c = (predicted == labels).squeeze()
for i in range(3):
    label = labels[i]
    class_correct[label] += c[i].item()
    class_total[label] += 1

for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

Accuracy of actinic keratoses :  0 %
Accuracy of basal cell carcinoma :  0 %
Accuracy of benign keratosis-like lesions : 43 %
Accuracy of dermatofibroma :  0 %
Accuracy of melanoma : 40 %
Accuracy of melanocytic nevi : 80 %
Accuracy of vascular lesions : 18 %

```

Confusion Matrix

```

confusion_matrix = torch.zeros(len(classes), len(classes))
with torch.no_grad():
    for data in test_data_loader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        for t, p in zip(labels.view(-1), predicted.view(-1)):
            confusion_matrix[t.long(), p.long()] += 1

cm = confusion_matrix.numpy()

```

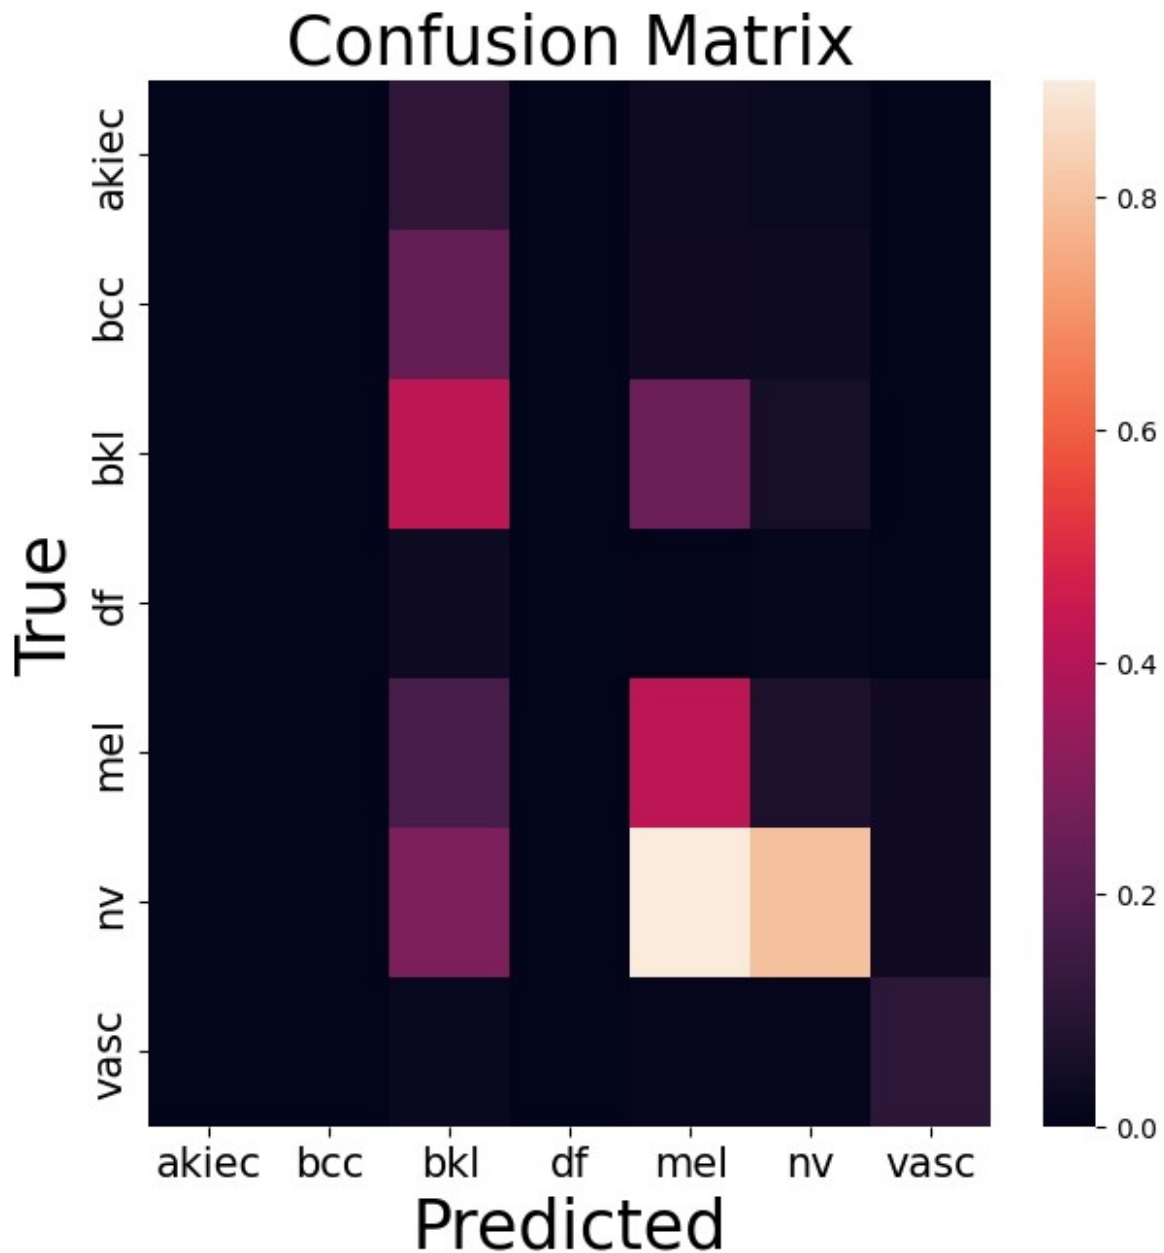
i have removed np. from (cm.astype(float)) and just used float

```

fig, ax = plt.subplots(figsize=(7, 7))
sns.heatmap(cm / (cm.astype(float).sum(axis=1) + 1e-9), annot=False,
            ax=ax)

# labels, title and ticks
ax.set_xlabel('Predicted', size=25);
ax.set_ylabel('True', size=25);
ax.set_title('Confusion Matrix', size=25);
ax.xaxis.set_ticklabels(['akiec', 'bcc', 'bkl', 'df', 'mel',
                        'nv', 'vasc'], size=15); \
ax.yaxis.set_ticklabels(['akiec', 'bcc', 'bkl', 'df', 'mel', 'nv', 'vasc'],
                        size=15);

```



Grad cam

```
from collections.abc import Sequence

class _BaseWrapper(object):
    """
    Please modify forward() and backward() according to your task.
    """

    def __init__(self, model):
        super(_BaseWrapper, self).__init__()
        self.device = next(model.parameters()).device
```

```

        self.model = model
        self.handlers = [] # a set of hook function handlers

    def _encode_one_hot(self, ids):
        one_hot = torch.zeros_like(self.logits).to(self.device)
        one_hot.scatter_(1, ids, 1.0)
        return one_hot

    def forward(self, image):
        """
        Simple classification
        """
        self.model.zero_grad()
        self.logits = self.model(image)
        self.probs = F.softmax(self.logits, dim=1)
        return self.probs.sort(dim=1, descending=True)

    def backward(self, ids):
        """
        Class-specific backpropagation
        Either way works:
        1. self.logits.backward(gradient=one_hot, retain_graph=True)
        2. (self.logits * one_hot).sum().backward(retain_graph=True)
        """

        one_hot = self._encode_one_hot(ids)
        self.logits.backward(gradient=one_hot, retain_graph=True)

    def generate(self):
        raise NotImplementedError

    def remove_hook(self):
        """
        Remove all the forward/backward hook functions
        """
        for handle in self.handlers:
            handle.remove()

class GradCAM(_BaseWrapper):
    """
    "Grad-CAM: Visual Explanations from Deep Networks via Gradient-
    based Localization"
    https://arxiv.org/pdf/1610.02391.pdf
    Look at Figure 2 on page 4
    """

    def __init__(self, model, candidate_layers=None):
        super(GradCAM, self).__init__(model)
        self.fmap_pool = OrderedDict()

```

```

self.grad_pool = OrderedDict()
self.candidate_layers = candidate_layers # list

def forward_hook(key):
    def forward_hook_(module, input, output):
        # Save featuremaps
        self.fmap_pool[key] = output.detach()

    return forward_hook_

def backward_hook(key):
    def backward_hook_(module, grad_in, grad_out):
        # Save the gradients correspond to the featuremaps
        self.grad_pool[key] = grad_out[0].detach()

    return backward_hook_

# If any candidates are not specified, the hook is registered
to all the layers.
for name, module in self.model.named_modules():
    if self.candidate_layers is None or name in
self.candidate_layers:

self.handlers.append(module.register_forward_hook(forward_hook(name)))

self.handlers.append(module.register_backward_hook(backward_hook(name)
))

def _find(self, pool, target_layer):
    if target_layer in pool.keys():
        return pool[target_layer]
    else:
        raise ValueError("Invalid layer name:
{}".format(target_layer))

def _compute_grad_weights(self, grads):
    return F.adaptive_avg_pool2d(grads, 1)

def forward(self, image):
    self.image_shape = image.shape[2:]
    return super(GradCAM, self).forward(image)

def generate(self, target_layer):
    fmaps = self._find(self.fmap_pool, target_layer)
    grads = self._find(self.grad_pool, target_layer)
    weights = self._compute_grad_weights(grads)

    gcam = torch.mul(fmaps, weights).sum(dim=1, keepdim=True)
    gcam = F.relu(gcam)

```

```

        gcam = F.interpolate(
            gcam, self.image_shape, mode="bilinear",
            align_corners=False
        )

        B, C, H, W = gcam.shape
        gcam = gcam.view(B, -1)
        gcam -= gcam.min(dim=1, keepdim=True)[0]
        gcam /= gcam.max(dim=1, keepdim=True)[0]
        gcam = gcam.view(B, C, H, W)

    return gcam

def demo2(image, label, model):
    """
    Generate Grad-CAM
    """
    # Model
    model = model
    model.to(device)
    model.eval()

    # The layers
    target_layers = ["conv2"]
    target_class = label

    # Images
    images = image.unsqueeze(0)
    gcam = GradCAM(model=model)
    probs, ids = gcam.forward(images)
    ids_ = torch.LongTensor([[target_class]] * len(images)).to(device)
    gcam.backward(ids=ids_)

    for target_layer in target_layers:
        print("Generating Grad-CAM @{}".format(target_layer))

        # Grad-CAM
        regions = gcam.generate(target_layer=target_layer)
        for j in range(len(images)):
            print(
                "\t#{}: {} ({:.5f})".format(
                    j, classes[target_class], float(probs[ids ==
target_class])
                )
            )

            gcam=regions[j, 0]
            plt.imshow(gcam.cpu())
            plt.show()


```

```
image, label = next(iter(test_data_loader))
# Load the model
model = net
# Grad cam
demo2(image[0].to(device), label[0].to(device), model)
```

```
image = np.transpose(image[0], (1,2,0))
image2 = np.add(np.multiply(image.numpy(),
np.array(norm_std)) ,np.array(norm_mean))
print("True Class: ", classes[label[0].cpu()])
plt.imshow(image)
plt.show()
plt.imshow(image2)
plt.show()
```



Generating Grad-CAM @conv2
#0: melanocytic nevi (0.61917)

```
P:\Python_root\Lib\site-packages\torch\nn\modules\module.py:1352:
UserWarning: Using a non-full backward hook when the forward contains
multiple autograd Nodes is deprecated and will be removed in future
versions. This hook will be missing some grad_input. Please use
register_full_backward_hook to get the documented behavior.
  warnings.warn("Using a non-full backward hook when the forward
contains multiple autograd Nodes ")
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

True Class: melanocytic nevi



Analysis of the results

As we can see from the results of the LeNet model, our system is not capable of processing the complexity of the given input images. Our final accuracy on the test data was 61%. About 39% of the images are missclassified, which is a terrible performance for any clinical use case.

These results could be substantially improved if we opt for a deeper, more complex network architecture than LeNet, which will allow for a richer learning of the corresponding image features.

Switching to superior network architecture:Resnet18

```
# Importing the neural network module from the PyTorch library.
from torch import nn
# Getting the number of classes.
num_classes = len(classes)
#Loading a pre-trained ResNet-18 model from torchvision
net = torchvision.models.resnet18(pretrained = True)

# We replace last layer of resnet to match our number of classes which
is 7
net.fc = nn.Linear(512, num_classes)
# Moving the model to the specified device(in this case GPU-cuda)
net = net.to(device)

/usr/local/lib/python3.10/dist-packages/torchvision/models/
_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
since 0.13 and may be removed in the future, please use 'weights'
instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:2
23: UserWarning: Arguments other than a weight enum or `None` for
'weights' are deprecated since 0.13 and may be removed in the future.
The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-
f37072fd.pth" to
/usr/cs/grad/masters/2024/akalapa/.cache/torch/hub/checkpoints/resnet
18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:01<00:00, 29.5MB/s]
```

1. Define a Loss function and Optimizer

Let's use a Classification Cross-Entropy loss.

$$H_{y'}(y) := - \sum_i y_i' \log(y_i)$$

The most common and effective Optimizer currently used is **Adam: Adaptive Moments**. You can look [here](#) for more information.

```
# Importing the optimizer module from PyTorch
import torch.optim as optim
# Moving the class weights to the specified device for computation
class_weights = class_weights.to(device)
# We Define the criterion for calculating the loss, which is cross-
entropy loss and Using class weights to handle class imbalance if any.
criterion = nn.CrossEntropyLoss(weight = class_weights)
# We are Using the Adam optimizer to update the parameters of the
network during training.
optimizer = optim.Adam(net.parameters(), lr=1e-5)
# Printing the architecture of the neural network.
print(net)
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2),
padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
)
```

```

    )
    (layer2): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2),
bias=False)
          (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (layer3): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2),
bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )
)

```

```

        (1): BasicBlock(
          (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (layer4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2),
bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Linear(in_features=512, out_features=7, bias=True)
  )

```

These are some helper functions to evaluate the training process.

```

# Importing accuracy_score function from sklearn.
from sklearn.metrics import accuracy_score
# Defining a function to calculate accuracy.
def get_accuracy(predicted, labels):
    ## Initializing variables to count batch length and correct predictions
    batch_len, correct= 0, 0
    ## Calculating the batch length
    batch_len = labels.size(0)
    # Counting the number of correct predictions.
    correct = (predicted == labels).sum().item()
    # Returning batch length and number of correct predictions.
    return batch_len, correct
# Defining a function to evaluate the model
def evaluate(model, val_loader):
    # Initializing variable to accumulate losses
    losses= 0
    # Initializing variable to count total number of samples.
    num_samples_total=0
    # Initializing variable to count total number of correct predictions.
    correct_total=0
    # Setting the model to evaluation mode.
    model.eval()
    # Iterating through the validation loader using for loop
    for inputs, labels in val_loader:
        # Moving inputs and labels to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)
        # Forward pass through the model.
        out = model(inputs)
        # Getting the index of the maximum value along the second dimension.
        _, predicted = torch.max(out, 1)
        # loss Calculation
        loss = criterion(out, labels)
        # Accumulating the loss.
        losses += loss.item()
        # Getting batch length and correct predictions.
        b_len, corr = get_accuracy(predicted, labels)
        # Accumulating the total number of samples.
        num_samples_total +=b_len
        # Accumulating the total number of correct predictions.
        correct_total +=corr
    #Accuracy calculation
    accuracy = correct_total/num_samples_total
    #Average loss calculation
    losses = losses/len(val_loader)
    #Returning average loss and accuracy.
    return losses, accuracy

```

Train the network

This is when things start to get interesting. We simply loop over the training data iterator, and feed the inputs to the network and optimize.

```
# number of loops over the dataset
num_epochs = 50
# Lists to store training and validation accuracy and losses for each
epoch.
accuracy = []
val_accuracy = []
losses = []
val_losses = []

# Looping over each epoch
for epoch in range(num_epochs):
    #Deining variables to store values.
    running_loss = 0.0
    correct_total= 0.0
    num_samples_total=0.0
    # Iterating over the training data loader
    for i, data in enumerate(train_data_loader):
        # get the inputs
        inputs, labels = data
        # Moving inputs and labels to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)
        # set the parameter gradients to zero
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        #compute accuracy
        # Extracting the predicted labels by selecting the class with
the highest probability.
        _, predicted = torch.max(outputs, 1)
        # Calculating the batch length (number of samples) and the
number of correct predictions.
        b_len, corr = get_accuracy(predicted, labels)
        # Accumulating the total number of samples processed in this
epoch.
        num_samples_total +=b_len
        # Accumulating the total number of correct predictions in this
epoch.
        correct_total +=corr
        # Accumulating the training loss for this batch.
```



```

        running_loss += loss.item()

    # Calculating average training loss and accuracy for the epoch.
    # Calculating the average training loss for this epoch.
    running_loss /= len(train_data_loader)
    # Calculating the training accuracy for this epoch.
    train_accuracy = correct_total/num_samples_total
    # Evaluating the model on the validation set to calculate
    validation loss and accuracy.
    val_loss, val_acc = evaluate(net, validation_data_loader)
    # Printing the current epoch.
    print('Epoch: %d' %(epoch+1))
    # Printing the average training loss and accuracy for the epoch.
    print('Loss: %.3f Accuracy: %.3f' %(running_loss, train_accuracy))
    # Printing the validation loss and accuracy for the epoch.
    print('Validation Loss: %.3f Val Accuracy: %.3f' %(val_loss,
    val_acc))
    # Appending the average training loss for this epoch to the list
    of losses.
    losses.append(running_loss)
    # Appending the validation loss for this epoch to the list of
    validation losses.
    val_losses.append(val_loss)
    # Appending the training accuracy for this epoch to the list of
    accuracies.
    accuracy.append(train_accuracy)
    # Appending the validation accuracy for this epoch to the list of
    validation accuracies.
    val_accuracy.append(val_acc)
## Printing a message indicating the end of training
print('Finished Training')

```

```

Epoch: 1
Loss: 1.791 Accuracy:0.137
Validation Loss: 1.558 Val Accuracy: 0.316
Epoch: 2
Loss: 1.124 Accuracy:0.605
Validation Loss: 1.086 Val Accuracy: 0.616
Epoch: 3
Loss: 0.911 Accuracy:0.662
Validation Loss: 0.920 Val Accuracy: 0.677
Epoch: 4
Loss: 0.762 Accuracy:0.700
Validation Loss: 1.020 Val Accuracy: 0.618
Epoch: 5
Loss: 0.705 Accuracy:0.709
Validation Loss: 0.948 Val Accuracy: 0.627
Epoch: 6
Loss: 0.658 Accuracy:0.718
Validation Loss: 0.862 Val Accuracy: 0.672

```

Epoch: 7
Loss: 0.602 Accuracy:0.739
Validation Loss: 0.816 Val Accuracy: 0.736
Epoch: 8
Loss: 0.575 Accuracy:0.756
Validation Loss: 0.816 Val Accuracy: 0.731
Epoch: 9
Loss: 0.548 Accuracy:0.751
Validation Loss: 0.834 Val Accuracy: 0.724
Epoch: 10
Loss: 0.548 Accuracy:0.762
Validation Loss: 0.801 Val Accuracy: 0.708
Epoch: 11
Loss: 0.518 Accuracy:0.772
Validation Loss: 0.776 Val Accuracy: 0.740
Epoch: 12
Loss: 0.465 Accuracy:0.784
Validation Loss: 0.853 Val Accuracy: 0.711
Epoch: 13
Loss: 0.453 Accuracy:0.787
Validation Loss: 0.798 Val Accuracy: 0.733
Epoch: 14
Loss: 0.411 Accuracy:0.793
Validation Loss: 0.828 Val Accuracy: 0.736
Epoch: 15
Loss: 0.407 Accuracy:0.806
Validation Loss: 0.803 Val Accuracy: 0.751
Epoch: 16
Loss: 0.409 Accuracy:0.807
Validation Loss: 0.882 Val Accuracy: 0.769
Epoch: 17
Loss: 0.375 Accuracy:0.812
Validation Loss: 0.792 Val Accuracy: 0.745
Epoch: 18
Loss: 0.348 Accuracy:0.822
Validation Loss: 0.792 Val Accuracy: 0.771
Epoch: 19
Loss: 0.352 Accuracy:0.822
Validation Loss: 0.798 Val Accuracy: 0.795
Epoch: 20
Loss: 0.306 Accuracy:0.829
Validation Loss: 0.918 Val Accuracy: 0.795
Epoch: 21
Loss: 0.333 Accuracy:0.830
Validation Loss: 0.778 Val Accuracy: 0.770
Epoch: 22
Loss: 0.304 Accuracy:0.833
Validation Loss: 0.720 Val Accuracy: 0.766
Epoch: 23

Loss: 0.280 Accuracy:0.843
Validation Loss: 0.913 Val Accuracy: 0.762
Epoch: 24
Loss: 0.292 Accuracy:0.848
Validation Loss: 0.904 Val Accuracy: 0.734
Epoch: 25
Loss: 0.269 Accuracy:0.845
Validation Loss: 0.795 Val Accuracy: 0.785
Epoch: 26
Loss: 0.246 Accuracy:0.855
Validation Loss: 0.837 Val Accuracy: 0.802
Epoch: 27
Loss: 0.224 Accuracy:0.863
Validation Loss: 0.821 Val Accuracy: 0.770
Epoch: 28
Loss: 0.244 Accuracy:0.863
Validation Loss: 0.838 Val Accuracy: 0.792
Epoch: 29
Loss: 0.209 Accuracy:0.878
Validation Loss: 0.853 Val Accuracy: 0.792
Epoch: 30
Loss: 0.259 Accuracy:0.856
Validation Loss: 0.842 Val Accuracy: 0.775
Epoch: 31
Loss: 0.211 Accuracy:0.874
Validation Loss: 0.776 Val Accuracy: 0.789
Epoch: 32
Loss: 0.198 Accuracy:0.881
Validation Loss: 0.871 Val Accuracy: 0.812
Epoch: 33
Loss: 0.193 Accuracy:0.880
Validation Loss: 0.847 Val Accuracy: 0.797
Epoch: 34
Loss: 0.178 Accuracy:0.891
Validation Loss: 0.872 Val Accuracy: 0.813
Epoch: 35
Loss: 0.180 Accuracy:0.888
Validation Loss: 0.861 Val Accuracy: 0.814
Epoch: 36
Loss: 0.183 Accuracy:0.888
Validation Loss: 0.844 Val Accuracy: 0.803
Epoch: 37
Loss: 0.151 Accuracy:0.899
Validation Loss: 0.880 Val Accuracy: 0.811
Epoch: 38
Loss: 0.155 Accuracy:0.900
Validation Loss: 0.740 Val Accuracy: 0.772
Epoch: 39
Loss: 0.165 Accuracy:0.897

```
Validation Loss: 0.859  Val Accuracy: 0.802
Epoch: 40
Loss: 0.132  Accuracy:0.907
Validation Loss: 0.929  Val Accuracy: 0.779
Epoch: 41
Loss: 0.140  Accuracy:0.908
Validation Loss: 0.913  Val Accuracy: 0.803
Epoch: 42
Loss: 0.149  Accuracy:0.903
Validation Loss: 1.051  Val Accuracy: 0.828
Epoch: 43
Loss: 0.173  Accuracy:0.899
Validation Loss: 1.040  Val Accuracy: 0.785
Epoch: 44
Loss: 0.156  Accuracy:0.904
Validation Loss: 0.940  Val Accuracy: 0.812
Epoch: 45
Loss: 0.120  Accuracy:0.922
Validation Loss: 0.962  Val Accuracy: 0.740
Epoch: 46
Loss: 0.112  Accuracy:0.919
Validation Loss: 1.047  Val Accuracy: 0.825
Epoch: 47
Loss: 0.122  Accuracy:0.924
Validation Loss: 0.942  Val Accuracy: 0.835
Epoch: 48
Loss: 0.113  Accuracy:0.921
Validation Loss: 1.019  Val Accuracy: 0.822
Epoch: 49
Loss: 0.137  Accuracy:0.913
Validation Loss: 0.988  Val Accuracy: 0.802
Epoch: 50
Loss: 0.108  Accuracy:0.927
Validation Loss: 1.005  Val Accuracy: 0.826
Finished Training
```


Plot the training and validation loss curves.

```
# Creating a plot of training and validation losses.
epoch = range(1, num_epochs + 1)
# Generating a list of epochs.
fig = plt.figure(figsize=(10, 15))
# Creating a new figure with a specific size.
plt.subplot(2, 1, 2)
# Creating a subplot with 2 rows, 1 column, and selecting the second subplot.
plt.plot(epoch, losses, label='Training loss')
# Plotting the training loss over epochs.
plt.plot(epoch, val_losses, label='Validation loss')
# Plotting the validation loss over epochs.
```

```
plt.title('Training and Validation Loss')
# Setting the title of the plot.
plt.xlabel('Epochs')
# Labeling the x-axis.
plt.legend()
# Adding a legend to the plot.
plt.figure()
# Creating a new figure (not necessary here, as we already have one).
plt.show()
# Displaying the plot.
# Creating a plot of training and validation accuracies.
fig = plt.figure(figsize=(10, 15))
# Creating a new figure with a specific size.
plt.subplot(2, 1, 2)
# Creating a subplot with 2 rows, 1 column, and selecting the second
subplot.
plt.plot(epoch, accuracy, label='Training accuracy')
# Plotting the training accuracy over epochs.
plt.plot(epoch, val_accuracy, label='Validation accuracy')
# Plotting the validation accuracy over epochs.
plt.title('Training and Validation Accuracy')
# Setting the title of the plot.
plt.xlabel('Epochs')
# Labeling the x-axis.
plt.legend()
# Adding a legend to the plot.
plt.figure()
# Creating a new figure (not necessary here, as we already have one).
plt.show()
# Displaying the plot.
```



<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>

Test the network on the test data

We have trained the network over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs, and checking it against the ground-truth. If the prediction is correct, we add the sample to the list of correct predictions.

Okay, first step. Let us display an image from the test set to get familiar.

```
fig = plt.figure(figsize=(10, 15))
# Creating a figure with a specific size.

# Iterate over the test data loader directly
for images, labels in test_data_loader:
    # Iterate over the test data loader directly
    imshow(torchvision.utils.make_grid(images))
    # Displaying the images in a grid.
```

```

    print('GroundTruth: ', ' '.join('%5s, ' % classes[labels[j]] for
j in range(len(labels))))
    # Printing the ground truth labels for the images.
    break
    # Break after the first batch to visualize it.


```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```

GroundTruth: melanocytic nevi, melanocytic nevi, melanocytic
nevi, melanocytic nevi, melanocytic nevi, melanocytic nevi,
melanocytic nevi, melanoma, melanocytic nevi, melanocytic nevi,

```



Okay, now let us check the performance on the test network:

```

correct = 0 # Initializing a variable to count the number of
correctly classified images.
total = 0 # Initializing a variable to count the total number of
images.

net.eval() # Setting the network to evaluation mode.

with torch.no_grad(): # Temporarily disabling gradient calculation
for efficiency.
    for data in test_data_loader: # Iterating over the test data
loader.
        images, labels = data # Extracting images and labels from the
current batch.
        images, labels = images.to(device), labels.to(device) #
Moving images and labels to the specified device (e.g., GPU).
        outputs = net(images) # Passing images through the network to
obtain outputs.
        _, predicted = torch.max(outputs.data, 1) # Getting the
predicted labels by selecting the class with the highest probability.
        total += labels.size(0) # Accumulating the total number of
images processed.
        correct += (predicted == labels).sum().item() # Accumulating
the number of correctly classified images.

```



```
print('Accuracy of the network on the test images: %d %%' % (
    100 * correct / total)) # Printing the accuracy of the network on
the test images.
```

Accuracy of the network on the test images: 83 %

```
# Temporarily disabling gradient calculation for efficiency.
with torch.no_grad():
    # Iterating over the test data loader.
    for data in test_data_loader:
        # Extracting images and labels from the current batch.
        images, labels = data
        # Moving images and labels to the specified device (e.g.,
GPU).
        images, labels = images.to(device), labels.to(device)
        # Passing images through the network to obtain outputs.
        outputs = net(images)
        # Getting the predicted labels by selecting the class with the
highest probability.
        _, predicted = torch.max(outputs, 1)
        # Creating a mask to identify correctly classified images.
        c = (predicted == labels).squeeze()
        # Iterating over a subset of images (e.g., first 3).
        for i in range(3):
            # Extracting the true label of the current image.
            label = labels[i]
            # Incrementing the count of correctly classified images
for the corresponding class.
            class_correct[label] += c[i].item()
            # Incrementing the total count of images for the
corresponding class.
            class_total[label] += 1

# Printing the accuracy for each class as a percentage
for i in range(len(classes)):
    # Iterating over each class.
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

Accuracy of actinic keratoses : 66 %
Accuracy of basal cell carcinoma : 84 %
Accuracy of benign keratosis-like lesions : 60 %
Accuracy of dermatofibroma : 71 %
Accuracy of melanoma : 54 %
Accuracy of melanocytic nevi : 91 %
Accuracy of vascular lesions : 99 %

```
# Temporarily disabling gradient calculation for efficiency.
with torch.no_grad():
```

```

# Iterating over the test data loader.
for data in test_data_loader:
    # Extracting images and labels from the current batch.
    images, labels = data
    # Moving images and labels to the specified device (e.g.,
GPU).
    images, labels = images.to(device), labels.to(device)
    # Passing images through the network to obtain outputs.
    outputs = net(images)
    # Getting the predicted labels by selecting the class with the
highest probability.
    _, predicted = torch.max(outputs, 1)
    # Creating a mask to identify correctly classified images.
    c = (predicted == labels).squeeze()
    # Iterating over a subset of images (e.g., first 3).
    for i in range(3):
        # Extracting the true label of the current image.
        label = labels[i]
        # Incrementing the count of correctly classified images
for the corresponding class.
        class_correct[label] += c[i].item()
        # Incrementing the total count of images for the
corresponding class.
        class_total[label] += 1

# Printing the accuracy for each class as a percentage
for i in range(len(classes)):
    # Iterating over each class.
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

Accuracy of actinic keratoses : 64 %
Accuracy of basal cell carcinoma : 90 %
Accuracy of benign keratosis-like lesions : 70 %
Accuracy of dermatofibroma : 49 %
Accuracy of melanoma : 47 %
Accuracy of melanocytic nevi : 90 %
Accuracy of vascular lesions : 99 %

```

Confusion Matrix

```

# Initializing a confusion matrix with zeros.
confusion_matrix = torch.zeros(len(classes), len(classes))

# Temporarily disabling gradient calculation for efficiency.
with torch.no_grad():
    # Iterating over the test data loader.
    for data in test_data_loader:
        # Extracting images and labels from the current batch.
        images, labels = data

```

```

    # Moving images and labels to the specified device (e.g.,
GPU).
    images, labels = images.to(device), labels.to(device)
    # Passing images through the network to obtain outputs.
    outputs = net(images)
    # Getting the predicted labels by selecting the class with the
highest probability.
    _, predicted = torch.max(outputs, 1)
    # Iterating over true and predicted labels.
    for t, p in zip(labels.view(-1), predicted.view(-1)):
        # Incrementing the corresponding entry in the confusion
matrix.
        confusion_matrix[t.long(), p.long()] += 1

# Printing the confusion matrix.
print(confusion_matrix)

# Converting the confusion matrix to a NumPy array for further
processing.
cm = confusion_matrix.numpy()

tensor([[4.7000e+01, 9.0000e+00, 4.0000e+00, 2.0000e+00, 0.0000e+00,
3.0000e+00,
        0.0000e+00],
        [8.0000e+00, 8.5000e+01, 5.0000e+00, 0.0000e+00, 1.0000e+00,
4.0000e+00,
        0.0000e+00],
        [1.6000e+01, 1.3000e+01, 1.4100e+02, 0.0000e+00, 1.6000e+01,
3.4000e+01,
        0.0000e+00],
        [0.0000e+00, 1.0000e+00, 2.0000e+00, 1.8000e+01, 0.0000e+00,
2.0000e+00,
        0.0000e+00],
        [1.2000e+01, 7.0000e+00, 1.3000e+01, 0.0000e+00, 1.3100e+02,
5.9000e+01,
        1.0000e+00],
        [4.0000e+00, 2.1000e+01, 4.0000e+01, 3.0000e+00, 5.4000e+01,
1.2170e+03,
        2.0000e+00],
        [0.0000e+00, 2.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00,
2.0000e+00,
        2.4000e+01]])

# Initializing a confusion matrix with zeros.
confusion_matrix = torch.zeros(len(classes), len(classes))

# Temporarily disabling gradient calculation for efficiency.
with torch.no_grad():
    # Iterating over the test data loader.
    for data in test_data_loader:

```


```

# Extracting images and labels from the current batch.
images, labels = data
# Moving images and labels to the specified device (e.g.,
GPU).
images, labels = images.to(device), labels.to(device)
# Passing images through the network to obtain outputs.
outputs = net(images)
# Getting the predicted labels by selecting the class with the
highest probability.
_, predicted = torch.max(outputs, 1)
# Iterating over true and predicted labels.
for t, p in zip(labels.view(-1), predicted.view(-1)):
    # Incrementing the corresponding entry in the confusion
matrix.
    confusion_matrix[t.long(), p.long()] += 1

# Printing the confusion matrix.
print(confusion_matrix)

# Converting the confusion matrix to a NumPy array for further
processing.
cm = confusion_matrix.numpy()

```



Grad cam

```
# Importing necessary libraries.
from collections.abc import Sequence

# Defining a base class for wrapper objects that modify model
behavior.
class _BaseWrapper(object):
    """
    Base class for wrapper objects that modify model behavior.
    """
```

```

# Constructor method to initialize the base wrapper.
def __init__(self, model):
    # Initialize the base class.
    super(_BaseWrapper, self).__init__()

    # Get the device of the model's parameters.
    self.device = next(model.parameters()).device

    # Set the model.
    self.model = model

    # Initialize a list to store hook function handlers.
    self.handlers = []

# Method to encode target class labels into one-hot vectors.
def _encode_one_hot(self, ids):
    """
    Encodes target class labels into one-hot vectors.
    """
    one_hot = torch.zeros_like(self.logits).to(self.device)
    one_hot.scatter_(1, ids, 1.0)
    return one_hot

# Method to perform forward pass through the model.
def forward(self, image):
    """
    Performs forward pass through the model.
    """
    # Zero gradients.
    self.model.zero_grad()

    # Forward pass.
    self.logits = self.model(image)
    self.probs = F.softmax(self.logits, dim=1)

    # Sort probabilities in descending order.
    return self.probs.sort(dim=1, descending=True)

# Method to perform backward pass to compute gradients.
def backward(self, ids):
    """
    Performs backward pass to compute gradients.
    """
    # Encode target labels into one-hot vectors.
    one_hot = self._encode_one_hot(ids)

    # Backward pass.
    self.logits.backward(gradient=one_hot, retain_graph=True)

# Method to generate the explanation.

```

```

def generate(self):
    """
    Generates the explanation.
    """
    raise NotImplementedError

# Method to remove all the forward/backward hook functions.
def remove_hook(self):
    """
    Removes all the forward/backward hook functions.
    """
    for handle in self.handlers:
        handle.remove()

# Defining a class for computing Grad-CAM explanations.
class GradCAM(_BaseWrapper):
    """
    Class for computing Grad-CAM explanations.
    """

    # Constructor method to initialize the GradCAM object.
    def __init__(self, model, candidate_layers=None):
        # Initialize the GradCAM object.
        super(GradCAM, self).__init__(model)

        # Initialize dictionaries to store feature maps and gradients.
        self.fmap_pool = OrderedDict()
        self.grad_pool = OrderedDict()

        # Store candidate layers for which hooks will be registered.
        self.candidate_layers = candidate_layers

        # Define hook functions for forward and backward passes.
        def forward_hook(key):
            def forward_hook_(module, input, output):
                # Save feature maps.
                self.fmap_pool[key] = output.detach()
            return forward_hook_

        def backward_hook(key):
            def backward_hook_(module, grad_in, grad_out):
                # Save gradients correspond to the feature maps.
                self.grad_pool[key] = grad_out[0].detach()
            return backward_hook_

        # Register hooks for the candidate layers.
        for name, module in self.model.named_modules():
            if self.candidate_layers is None or name in
self.candidate_layers:

```

```

self.handlers.append(module.register_forward_hook(forward_hook(name)))

self.handlers.append(module.register_backward_hook(backward_hook(name)))

    # Helper method to retrieve feature maps or gradients for a given layer.
    def _find(self, pool, target_layer):
        """
        Helper function to retrieve feature maps or gradients for a
        given layer.
        """
        if target_layer in pool.keys():
            return pool[target_layer]
        else:
            raise ValueError("Invalid layer name:
{}".format(target_layer))

    # Method to compute gradient weights using global average pooling.
    def _compute_grad_weights(self, grads):
        """
        Computes gradient weights using global average pooling.
        """
        return F.adaptive_avg_pool2d(grads, 1)

    # Method to perform forward pass and stores the image shape.
    def forward(self, image):
        """
        Performs forward pass and stores the image shape.
        """
        self.image_shape = image.shape[2:]
        return super(GradCAM, self).forward(image)

    # Method to generate Grad-CAM explanation for a specific target layer.
    def generate(self, target_layer):
        """
        Generates Grad-CAM explanation for a specific target layer.
        """
        # Retrieve feature maps and gradients for the target layer.
        fmaps = self._find(self.fmap_pool, target_layer)
        grads = self._find(self.grad_pool, target_layer)

        # Compute gradient weights.
        weights = self._compute_grad_weights(grads)

        # Generate Grad-CAM.
        gcam = torch.mul(fmaps, weights).sum(dim=1, keepdim=True)
        gcam = F.relu(gcam)

```



```

        gcam = F.interpolate(gcam, self.image_shape, mode="bilinear",
align_corners=False)

        # Normalize the Grad-CAM.
        B, C, H, W = gcam.shape
        gcam = gcam.view(B, -1)
        gcam -= gcam.min(dim=1, keepdim=True)[0]
        gcam /= gcam.max(dim=1, keepdim=True)[0]
        gcam = gcam.view(B, C, H, W)

        return gcam

def demo2(image, label, model):
    """
    Generates Grad-CAM for a given image and label using the specified
    model.
    """
    # Move the model to the appropriate device and set to evaluation
    mode.
    model = model
    model.to(device)
    model.eval()

    # Specify the target layers for Grad-CAM computation.
    target_layers = ["layer4"]
    target_class = label

    # Prepare the input image.
    images = image.unsqueeze(0)

    # Initialize GradCAM object.
    gcam = GradCAM(model=model)

    # Forward pass to obtain probabilities and predicted class IDs.
    probs, ids = gcam.forward(images)

    # Prepare target class IDs for backpropagation.
    ids_ = torch.LongTensor([[target_class]] * len(images)).to(device)

    # Perform backpropagation to compute gradients.
    gcam.backward(ids=ids_)

    # Generate Grad-CAM for each target layer.
    for target_layer in target_layers:
        print("Generating Grad-CAM @{}".format(target_layer))

        # Compute Grad-CAM regions.
        regions = gcam.generate(target_layer=target_layer)

        # Display Grad-CAM regions for each image.

```

```

        for j in range(len(images)):
            print(
                "\t#{:}: {} ({:.5f})".format(
                    j, classes[target_class], float(probs[ids ==
target_class])
                )
            )

            # Display the Grad-CAM region overlaid on the original
image.
            gcam_image = regions[j, 0]
            plt.imshow(gcam_image.cpu())
            plt.show()

```

```

# Get an image and label from the test dataset.
image, label = next(iter(test_data_loader))

```

```

# Load the model.
model = net

```

```

# Generate Grad-CAM for the image.
demo2(image[0].to(device), label[0].to(device), model)

```

```

# Display the original image and its normalized version.
image = np.transpose(image[0], (1,2,0))
image2 = np.add(np.multiply(image.numpy(), np.array(norm_std)),
np.array(norm_mean))
print("True Class: ", classes[label[0].cpu()])

```

```

# Display the original image.
plt.imshow(image)
plt.show()

```

```

# Display the normalized image.
plt.imshow(image2)
plt.show()

```

```


/usr/local/lib/python3.10/dist-packages/torch/nn/modules/
module.py:1359: UserWarning: Using a non-full backward hook when the
forward contains multiple autograd Nodes is deprecated and will be
removed in future versions. This hook will be missing some grad_input.
Please use register_full_backward_hook to get the documented behavior.
  warnings.warn("Using a non-full backward hook when the forward
contains multiple autograd Nodes ")

```

```



Generating Grad-CAM @layer4
#0: melanoma (0.98846)

```



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

True Class: melanoma



Conclusion

Training a neural network can be a daunting task, especially for a beginner. Here, are some useful practices to get the best out of your network.

- Training Ensembles — Combine learning from multiple networks.
- Always go for a lower learning rate.
- In cases of limited data try better augmentation techniques[20].
- Network architectures that have the appropriate depth for our problem — too many hyperparameters could lead to suboptimal results if we don't have enough images.
- Improving loss function and class balancing.

In this tutorial we learned how to train a deep neural network for the challenging task of skin-lesion classification. We experimented with two network architectures and provided insights in the attention of the models. Additionally, we achieved 83% overall accuracy on HAM10000 and provided you with more tips and tricks to tackle overfitting and class imbalance.

Now you have all the tools to not only beat our performance and participate in the exciting MICCAI Challenges, but to also solve many more medical imaging problems.

Happy training!